



# LINUX PROGRAMMING

---

## Inter-Process Communication (IPC)

# Outline

- Introduction to Inter-Process Communication
  - Inter-Process Communication (IPC) models
    - IPC in Shared-Memory Systems
      - Mapped Memory
      - Pipes
    - IPC in Message-Passing Systems
  - Examples of IPC Systems
  - Communication in Client-Server Systems
    - Sockets
- Labs



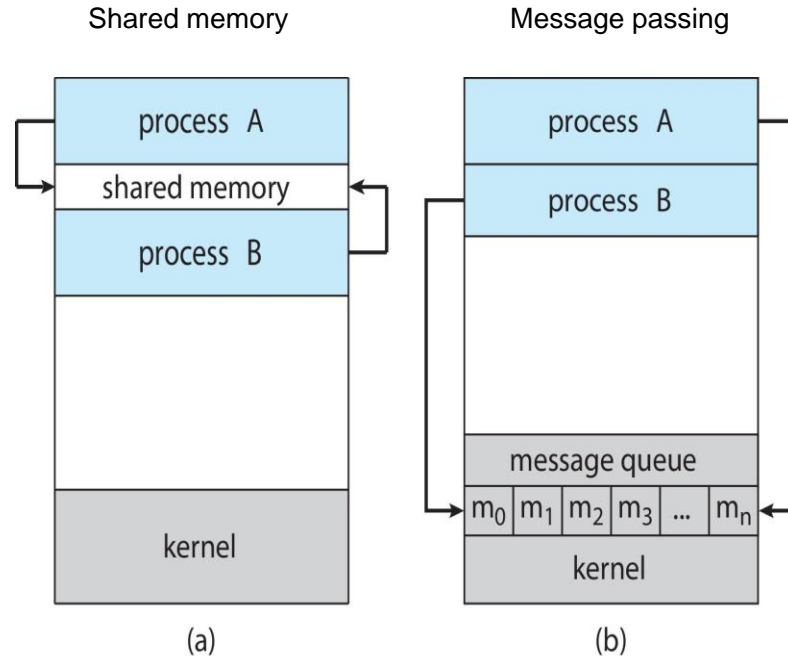
# Inter-Process Communication (IPC)

- Processes within a system may be *independent* or *cooperating*
  - *Independent process* does not share data with any other processes executing in the system
  - *Cooperating process* can affect or be affected by other processes, including sharing data
- Reasons for cooperating processes:
  - Information sharing
  - Computation speed-up
  - Modularity
  - Convenience
- Cooperating processes need *Inter-Process Communication (IPC)*



# Communication Models

- Types of IPC differ by the following criteria
  - Related/unrelated processes
  - Read/write operations
  - Number of processes permitted
  - Synchronization
- Two models of IPC
  - *Shared Memory*
  - *Message Passing*



# Shared Memory

# Inter-Process Communication – Shared Memory

- Shared memory: an *area of memory* shared among the processes that wish to communicate
  - allowing two or more processes to access the same memory as if they all called `malloc()` and were returned *pointers* to the same actual memory.
  - When one process changes the memory, all the other processes see the modification
- The communication is *under the control of the user processes*, not the operating system (kernel).
- Major issue is to provide a mechanism that will allow the user processes to *synchronize* their actions when they access shared memory (e.g., avoid race conditions).



# Synchronization: Producer-Consumer Problem

- Paradigm for cooperating processes: Producer-Consumer problem
  - *Producer process* produces information that will be consumed by a *consumer process*
- Buffer
  - *unbounded-buffer* places no practical limit on the size of the buffer
  - *bounded-buffer* assumes that there is a fixed buffer size



# Bounded-Buffer – Shared-Memory Solution

- Shared data

```
#define BUFFER_SIZE 10
```

```
typedef struct {
```

```
    ...
```

```
} item;
```

```
item buffer[BUFFER_SIZE];
```


```
int in = 0;
```

```
int out = 0;
```


- Solution is correct, but can only use `BUFFER_SIZE-1` elements

## Bounded Buffer

Block on:	Unblock on:
Producer: insert in full buffer	Consumer: item inserted
Consumer: remove from empty buffer	Producer: item removed



(a)



(b)





# Producer Process – Shared Memory

```
item next_produced;

while (true) {
    /* produce an item in next produced */
    while (((in + 1) % BUFFER_SIZE) == out)
        ; /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}
```



# Consumer Process – Shared Memory

```
item next_consumed;

while (true) {
    while (in == out)
        ; /* do nothing */
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;

    /* consume the item in next consumed */
}
```



# Inter-Process Communication - Mapped Memory

- Mapped memory permits different processes to communicate *via a shared file* (or as an easy way to access the contents of a file).
- Mapped memory forms an *association* between a file and a process's memory.
  - Linux splits the file into *page-sized chunks* and then copies them into *virtual memory pages* available in a process's address space.
- Mapped memory
  - allocating a buffer to hold a file's entire contents,
  - reading the file into the buffer,
  - writing the buffer back out to the file.



# Mapped Memory - Mapping an Ordinary File

- To map an ordinary file to a process's memory, use the `mmap` call.
  - first argument is the address,
  - second argument is the length of the map in bytes,
  - third argument specifies the protection on the mapped address range,
  - fourth argument is a flag value,
  - fifth argument is a file descriptor,
  - last argument is the offset from the beginning of the file (where starting the map).
  - If the call succeeds, it returns a pointer to the beginning of the memory.
- To finish with a memory mapping, release it by using `munmap` call.



# Mapped Memory - Shared Access to a File

- Different processes can communicate *using memory-mapped regions* associated with the same file.
- Specify the `MAP_SHARED` flag so that any writes to these regions are immediately transferred to the underlying file and made visible to other processes.
- Alternatively, you can force Linux to *incorporate buffered writes* into the disk file by calling `msync`.
- Users of memory-mapped regions must establish and follow a protocol to *avoid race conditions*.



# Message Passing

# Inter-Process Communication – Message Passing

- Mechanism for processes to *communicate* and to *synchronize* their actions
- Message system – processes communicate with each other *without resorting to shared variables*
- IPC facility provides two operations:
  - send(message)
  - receive(message)
- The message size is either *fixed* or *variable*



# Message Passing (Cont.)

- If processes **P** and **Q** wish to communicate, they need to:
  - establish a *communication link* between them
  - *exchange messages* via send/receive primitives
- Implementation issues:
  - How are links established?
  - Can a link be associated with more than two processes?
  - How many links can there be between every pair of communicating processes?
  - What is the capacity of a link?
  - Is the size of a message that the link can accommodate fixed or variable?
  - Is a link unidirectional or bi-directional?





# Message Passing (Cont.)

- Implementation of the communication link
  - *Physical*
    - Shared memory
    - Hardware bus
    - Network equipment
  - *Logical*
    - Direct or indirect
    - Synchronous or asynchronous
    - Automatic or explicit buffering
    - Network protocols



# Direct Communication

- Processes must *name* each other *explicitly*:
  - `send (P,message)` – send a message to process **P**
  - `receive (Q,message)` – receive a message from process **Q**
- Properties of the communication link
  - Links are established *automatically*
  - A link is associated with exactly *one pair* of communicating processes
  - Between each pair there exists *exactly one link*
  - The link may be *unidirectional*, but is usually *bi-directional*



# Indirect Communication

- Messages are *directed* and *received* from mailboxes (also referred to as *ports*)
  - Each mailbox has a *unique ID*
  - Processes can communicate only if they *share a mailbox*
- Properties of the communication link
  - Link established only if processes share a common mailbox
  - A link may be associated with many processes
  - Each pair of processes may share several communication links
  - Link may be unidirectional or bi-directional



# Indirect Communication (Cont.)

- Operations
  - *create* a new mailbox (or port)
  - *send* and *receive* messages through the mailbox
  - *destroy* a mailbox
- Primitives are defined as:
  - *send(A, message)* – send a message to mailbox A
  - *receive(A, message)* – receive a message from mailbox A



# Indirect Communication (Cont.)

- *Mailbox sharing*
  - Example
    - **P1**, **P2**, and **P3** share mailbox **A**,
    - **P1** sends; **P2** and **P3** receive.
    - *Who gets the message?*
  - Solutions
    - Allow a link to be associated with at most two processes
    - Allow only one process at a time to execute a receive operation
    - Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was



# Message Passing – Synchronization

- Message passing may be either *blocking* or *non-blocking*
- **Blocking** is considered as synchronous
  - *Blocking send* – the sender is blocked until the message is received
  - *Blocking receive* – the receiver is blocked until a message is available
- **Non-blocking** is considered as asynchronous
  - *Non-blocking send* – the sender sends the message and continue
  - *Non-blocking receive* – the receiver receives: a valid message, or Null message
- Different combinations possible
  - If both send and receive are blocking, we have a *rendezvous mechanism*



# Producer – Message Passing

```
message next_produced;
```

```
while (true) {  
    /* produce an item in next_produced */
```

```
    send(next_produced);
```

```
}
```



# Consumer – Message Passing

```
message next_consumed;
```

```
while (true) {  
    receive(next_consumed)
```

```
    /* consume the item in next_consumed */
```

```
}
```





# Buffering

- *Queue of messages* attached to the link.
- Implemented in one of three ways
  - *Zero capacity* – no messages are queued on a link
    - Sender must wait for the receiver (rendezvous)
  - *Bounded capacity* – finite length of n messages
    - Sender must wait if link is full
  - *Unbounded capacity* – infinite length
    - Sender never waits



# Examples of IPC Systems - POSIX

- *POSIX Shared Memory*
  - Process first creates shared memory segment
    - `shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);`
  - Also used to open an existing segment
  - Set the size of the object
    - `ftruncate(shm_fd, 4096);`
  - Use `mmap()` to memory-map a file pointer to the shared memory object
  - Reading and writing to shared memory are done by using the pointer returned by `mmap()`.



# IPC POSIX Producer – Consumer

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* strings written to shared memory */
    const char *message_0 = "Hello";
    const char *message_1 = "World!";

    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* create the shared memory object */
    shm_fd = shm.open(name, O_CREAT | O_RDWR, 0666);

    /* configure the size of the shared memory object */
    ftruncate(shm_fd, SIZE);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_WRITE, MAP_SHARED, shm_fd, 0);

    /* write to the shared memory object */
    sprintf(ptr, "%s", message_0);
    ptr += strlen(message_0);
    sprintf(ptr, "%s", message_1);
    ptr += strlen(message_1);

    return 0;
}
```

```
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* open the shared memory object */
    shm_fd = shm.open(name, O_RDONLY, 0666);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_READ, MAP_SHARED, shm_fd, 0);

    /* read from the shared memory object */
    printf("%s", (char *)ptr);

    /* remove the shared memory object */
    shm.unlink(name);

    return 0;
}
```



# Examples of IPC Systems - Mach

- Mach communication is *message-based*
  - Even system calls are *messages*
  - Each task gets two *ports* at creation – *Task Self port* and *Notify port*
  - Messages are sent and received using the `mach_msg()` function
  - Ports needed for communication, created via `mach_port_allocate()`
  - Send and receive are flexible, for example four options if mailbox full:
    - Wait indefinitely
    - Wait at most n milliseconds
    - Return immediately
    - Temporarily cache a message



# Mach Messages

```
#include<mach/mach.h>

struct message {
    mach_msg_header_t header;
    int data;
};

mach port t client;
mach port t server;
```



# Mach Message Passing - Client

```
/* Client Code */

struct message message;

// construct the header
message.header.msgh_size = sizeof(message);
message.header.msgh_remote_port = server;
message.header.msgh_local_port = client;

// send the message
mach_msg(&message.header, // message header
        MACH_SEND_MSG, // sending a message
        sizeof(message), // size of message sent
        0, // maximum size of received message - unnecessary
        MACH_PORT_NULL, // name of receive port - unnecessary
        MACH_MSG_TIMEOUT_NONE, // no time outs
        MACH_PORT_NULL // no notify port
    );
```



# Mach Message Passing - Server

```
/* Server Code */

struct message message;

// receive the message
mach_msg(&message.header, // message header
        MACH_RCV_MSG, // sending a message
        0, // size of message sent
        sizeof(message), // maximum size of received message
        server, // name of receive port
        MACH_MSG_TIMEOUT_NONE, // no time outs
        MACH_PORT_NULL // no notify port
);
```



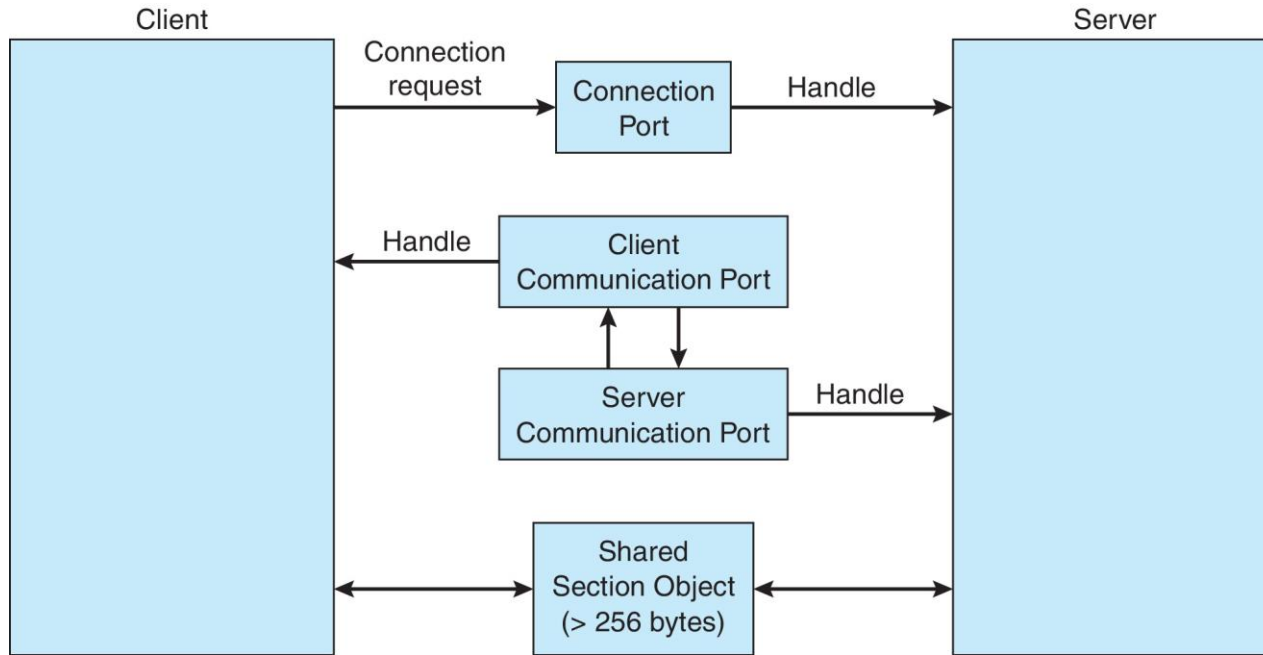
# Examples of IPC Systems – Windows

- Message-passing centric via advanced *Local Procedure Call (LPC)* facility
  - Only works between *processes on the same system*
  - Uses *ports* (like mailboxes) to establish and maintain communication channels
  - Communication works as follows:
    - The client opens a handle to the *subsystem's connection port* object
    - The client sends a *connection request*
    - The server creates *two private communication ports* and returns the handle to one of them to the client
    - The client and server use the corresponding port handle to *send messages* or callbacks and to *listen for replies*





# Local Procedure Calls in Windows



# Pipes


# Pipes

- Acts as a conduit allowing *2 processes* (or *threads*) to communicate
- Issues:
  - Is communication *unidirectional* or *bidirectional*?
  - In the case of two-way communication, is it *half* or *full-duplex*?
  - Must there exist a *relationship* (e.g., parent-child) between communicating processes?
  - Can the *pipes* be used over a network?
- *Ordinary pipes* – cannot be accessed from outside the process that created it. Typically, a parent process creates a pipe and uses it to communicate with a child process that it created.
- *Named pipes* – can be accessed without a parent-child relationship.



# Ordinary Pipes

- *Ordinary Pipes* allow communication in standard producer-consumer style
  - Producer writes to one end (the write-end of the pipe)
  - Consumer reads from the other end (the read-end of the pipe)
- Ordinary pipes are therefore *unidirectional*
- Require *parent-child relationship* between communicating processes



# Named Pipes

- *Named pipes* are more powerful than ordinary pipes
- Communication is *bidirectional*
- *No parent-child relationship* is necessary between the communicating processes
- *Several processes* can use the named pipe for communication
- Provided on both **UNIX** and **Windows** systems



# Creating Pipes

- In a shell, the symbol `|` creates a pipe.
- To create a pipe, invoke the `pipe` command.
  - Supply an integer array of size 2.
  - The call to `pipe` stores the reading file descriptor in array position 0 and the writing file descriptor in position 1.
  - Data written to the file descriptor `read_fd` can be read back from `write_fd`.
- Communication Between Parent and Child Processes
  - A call to `pipe` creates file descriptors, which are valid only within that process and its children.
  - When the process calls `fork`, file descriptors are copied to the new child process.
  - Pipes can connect only related processes.



# Ordinary Pipes – Redirecting the Standard Input, Output, and Error Streams

- Frequently, the need is to create a child process and set up one end of a pipe as its *standard input* or *standard output*.
  - Using the `dup2` call, you can equate one file descriptor with another.
- A common use of pipes is to send data to or receive data from a *program* being run in a *subprocess*.
  - The `popen` and `pclose` functions ease this paradigm by eliminating the need to invoke `pipe`, `fork`, `dup2`, `exec`, and `fdopen`.



# Named Pipe – First-In, First-Out (FIFO)

- A FIFO file is a pipe that has a name in the filesystem.
  - FIFOs are also called *named pipes*.
  - Processes on either end of the pipe need not be related to each other.
- *Creating a FIFO*: programmatically, using the `mkfifo` command.
- *Accessing a FIFO*: like an ordinary file.
  - To communicate through a FIFO, one program must open it for *writing*, and another program must open it for *reading*.
- Differences from *Windows Named Pipes*
  - Main differences: named pipes for Win32 function more like *sockets*.
  - *Win32 named pipes* can connect processes on separate computers connected *via a network*.





# Communications in Client-Server Systems

- *Sockets*
  - A *socket* is defined as an *endpoint* for communication
  - It is a concatenation of *IP address* and *port* – a number included at start of message packet to differentiate network services on a host
    - E.g., The socket 161.25.19.8:1625 refers to port 1625 on host 161.25.19.8
  - *Communication* consists of *a pair of sockets*
  - All ports below 1024 are *well-known*, used for *standard services*
  - Special IP address *127.0.0.1* (*loopback*) to refer to the system on which process is running
- *Remote Procedure Calls (RPC)*




# Socket Concepts

- Creating a socket must specify three parameters:
  - *Communication style*: how the socket treats transmitted data and specifies the number of communication partners? connection-oriented, connectionless (or datagram),
  - *Namespace*: how socket addresses are written? in the "local namespace" or in the "Internet namespace",
  - *Protocol*: how data is transmitted? **TCP** or **UDP**.
- System Calls
  - `Socket` — Creates a socket
  - `Closes` — Destroys a socket
  - `Connect` — Creates a connection between two sockets
  - `Bind` — Labels a server socket with an address
  - `Listen` — Configures a socket to accept conditions
  - `Accept` — Accepts a connection and creates a new socket for the connection



# Socket Communication



# Example app: TCP Server in Python

## *Python TCPServer*

	<pre>from socket import *</pre>
	<pre>serverPort = 12000</pre>
create TCP welcoming socket →	<pre>serverSocket = socket(AF_INET, SOCK_STREAM)</pre>
	<pre>serverSocket.bind(('', serverPort))</pre>
server begins listening for incoming TCP requests →	<pre>serverSocket.listen(1)</pre>
	<pre>print 'The server is ready to receive'</pre>
loop forever →	<pre>while True:</pre>
server waits on accept() for incoming requests, new socket created on return →	<pre>connectionSocket, addr = serverSocket.accept()</pre>
	<pre>    sentence = connectionSocket.recv(1024).decode()</pre>
read bytes from socket (but not address as in UDP) →	<pre>    capitalizedSentence = sentence.upper()</pre>
	<pre>    connectionSocket.send(capitalizedSentence.encode())</pre>
close connection to this client (but <i>not</i> welcoming socket) →	<pre>    connectionSocket.close()</pre>



# Sockets in Java – Server

- Three types of sockets
  - Connection-oriented (TCP)
  - Connectionless (UDP)
  - MulticastSocket class– data can be sent to multiple recipients
- Consider this “Date” server in Java:

```
import java.net.*;
import java.io.*;

public class DateServer
{
    public static void main(String[] args) {
        try {
            ServerSocket sock = new ServerSocket(6013);

            /* now listen for connections */
            while (true) {
                Socket client = sock.accept();

                PrintWriter pout = new
                    PrintWriter(client.getOutputStream(), true);

                /* write the Date to the socket */
                pout.println(new java.util.Date().toString());

                /* close the socket and resume */
                /* listening for connections */
                client.close();
            }
        }
        catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```



# Example app: TCP Client in Python

## *Python TCPClient*

create TCP socket for server,  
remote port 12000

No need to attach server name, port

```
from socket import *
serverName = 'servername'
serverPort = 12000
clientSocket = socket(AF_INET, SOCK_STREAM)
clientSocket.connect((serverName, serverPort))
sentence = raw_input('Input lowercase sentence:')
clientSocket.send(sentence.encode())
modifiedSentence = clientSocket.recv(1024)
print ('From Server:', modifiedSentence.decode())
clientSocket.close()
```



# Sockets in Java – Client

- The equivalent “Date” client

```
import java.net.*;
import java.io.*;

public class DateClient
{
    public static void main(String[] args) {
        try {
            /* make connection to server socket */
            Socket sock = new Socket("127.0.0.1",6013);

            InputStream in = sock.getInputStream();
            BufferedReader bin = new
                BufferedReader(new InputStreamReader(in));

            /* read the date from the socket */
            String line;
            while ( (line = bin.readLine()) != null)
                System.out.println(line);

            /* close the socket connection*/
            sock.close();
        }
        catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```



# Remote Procedure Calls (RPC)



# Remote Procedure Calls

- *Remote Procedure Call (RPC)* abstracts procedure calls between processes *on networked systems*
  - Again, uses *ports* for service differentiation
- *Stubs* – proxies for the actual procedure on the server and client sides
  - The *client-side stub* locates the server and marshals the parameters
  - The *server-side stub* receives this message, unpacks the marshalled parameters, and performs the procedure on the server
- On **Windows**, stub code compiled from specifications written in *Microsoft Interface Definition Language (MIDL)*

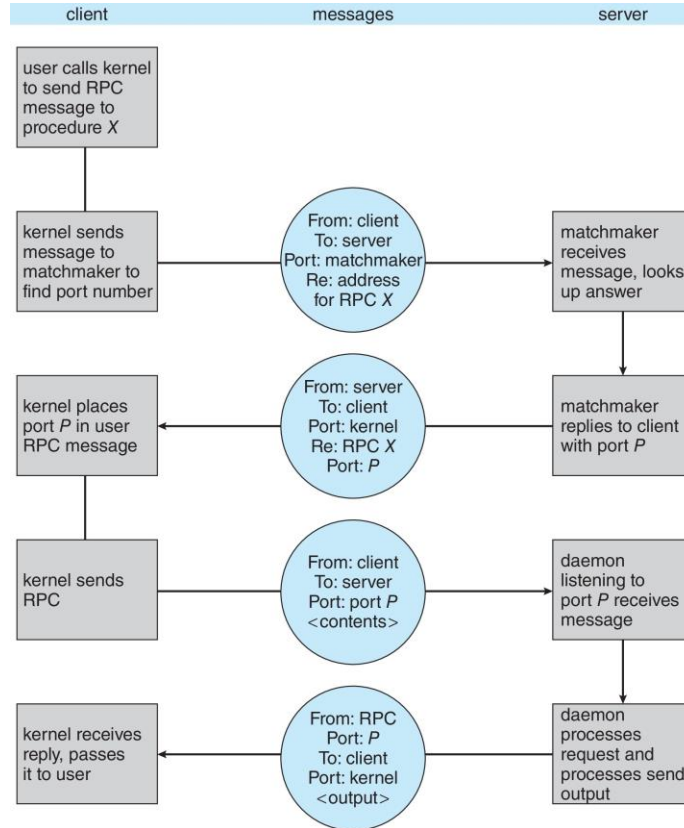


# Remote Procedure Calls (Cont.)

- Data representation handled via *External Data Representation (XDR)* format to account for different architectures
  - E.g., Big-endian (Motorola) and little-endian (Intel x86)
- *Remote communication* has more failure scenarios than local
  - Messages can be delivered exactly once rather than at most once
- OS typically provides a *rendezvous* (or *matchmaker*) service to connect Client and Server



# Execution of RPC



# Summary (Cont.)

- When *shared memory* is used for communication between processes, two (or more) processes share the same region of memory. **POSIX** provides an API for shared memory.
- Two processes may communicate by *exchanging messages* with one another using message passing. The **Mach** operating system uses message passing as its primary form of inter-process communication. **Windows** provides a form of message passing as well.



# Summary (Cont.)

- A *pipe* provides a conduit for two processes to communicate. There are two forms of pipes, *ordinary* and *named*. Ordinary pipes are designed for communication between processes that have a parent-child relationship. Named pipes are more general and allow several processes to communicate.
- **UNIX** systems provide ordinary pipes through the `pipe()` system call. *Ordinary pipes* have *a read end* and *a write end*. A parent process can, for example, send data to the pipe using its write end, and the child process can read it from its read end. *Named pipes* in UNIX are termed FIFOs.



# Summary (Cont.)

- **Windows** systems also provide two forms of pipes—anonymous and named pipes. *Anonymous pipes* are similar to UNIX ordinary pipes. They are unidirectional and employ parent-child relationships between the communicating processes. *Named pipes* offer a richer form of inter-process communication than the UNIX counterpart, FIFOs.
- Two common forms of client-server communication are *sockets* and *remote procedure calls (RPCs)*. Sockets allow two processes on different machines to communicate *over a network*. RPCs (used in Android) abstract the concept of function (procedure) calls in such a way that a function can be invoked on another process that may reside on a separate computer.





# THANK YOU !

---

## Center Of Computer Engineering