



LINUX PROGRAMMING

Processes

Outline

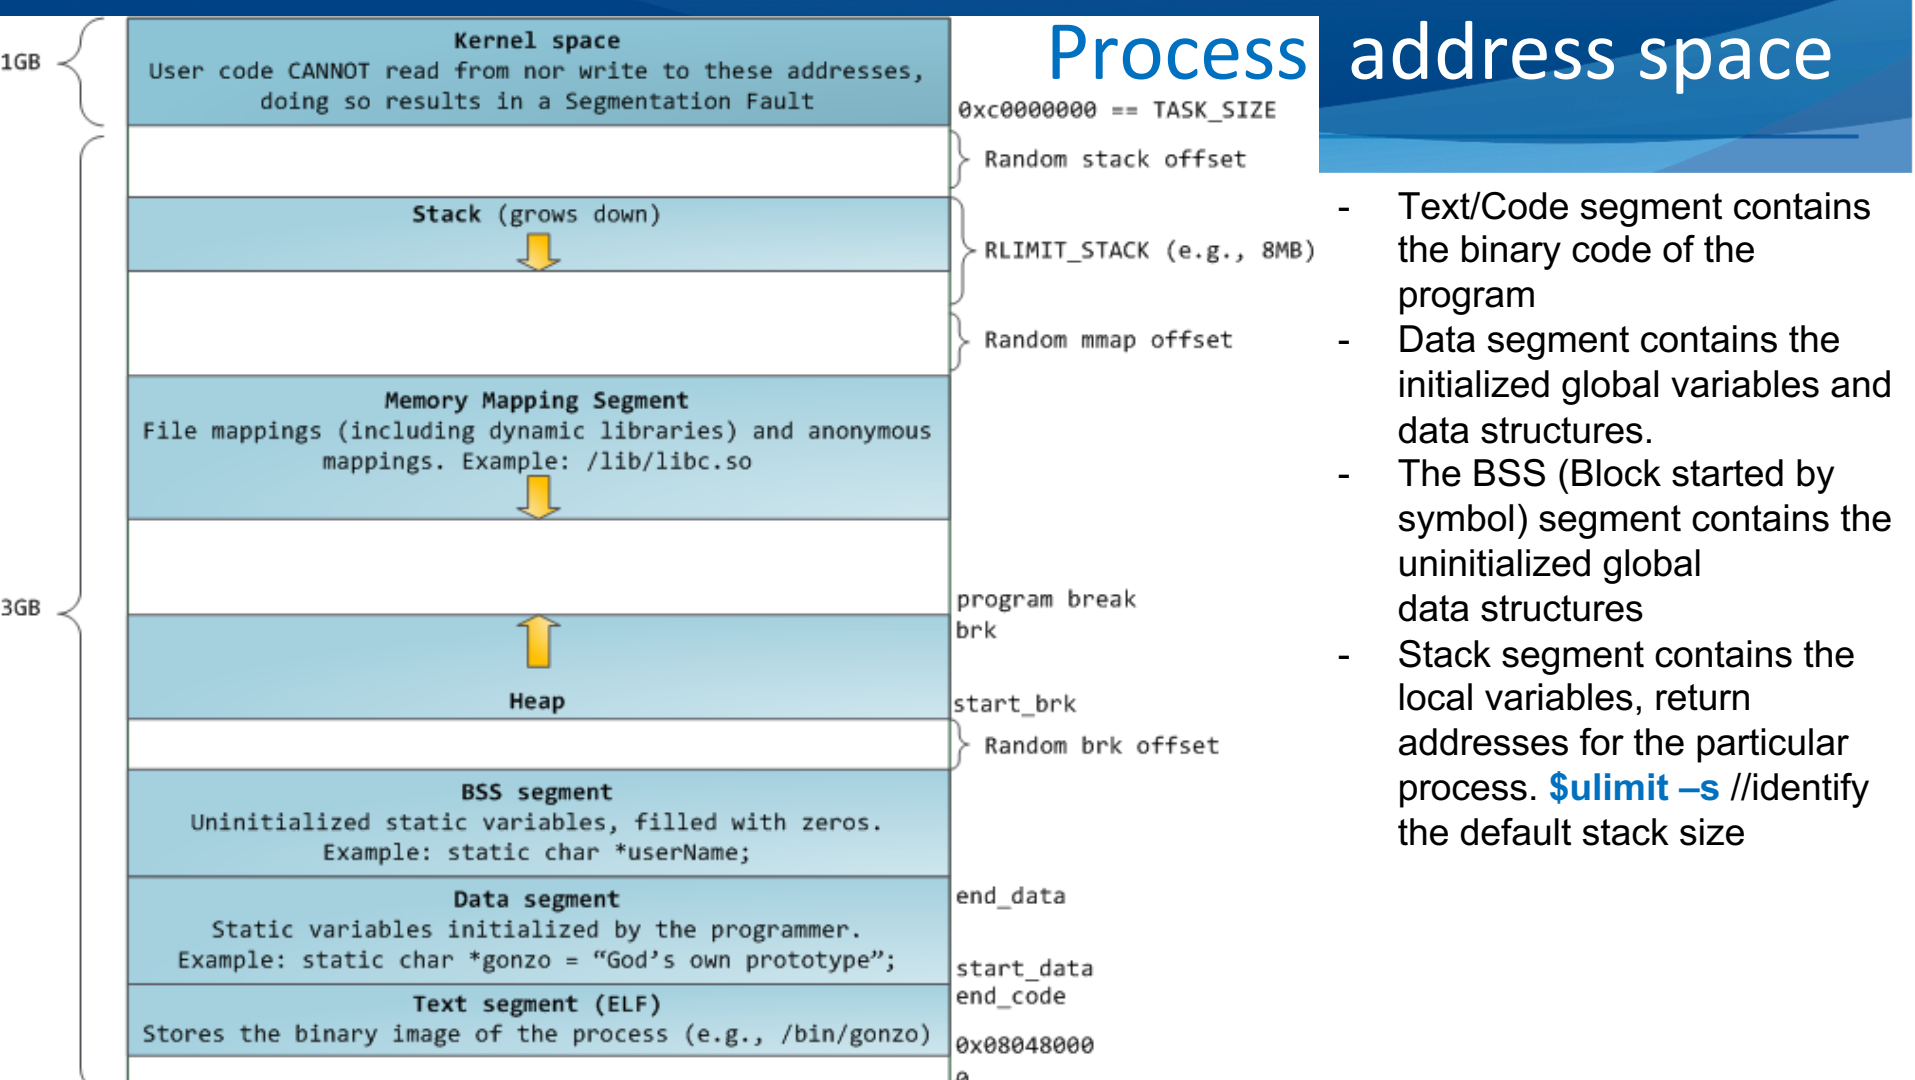
- Process Concept
- Process address space
- Process Scheduling
- Process creation
- Process termination
- Signal handling

Process Concept

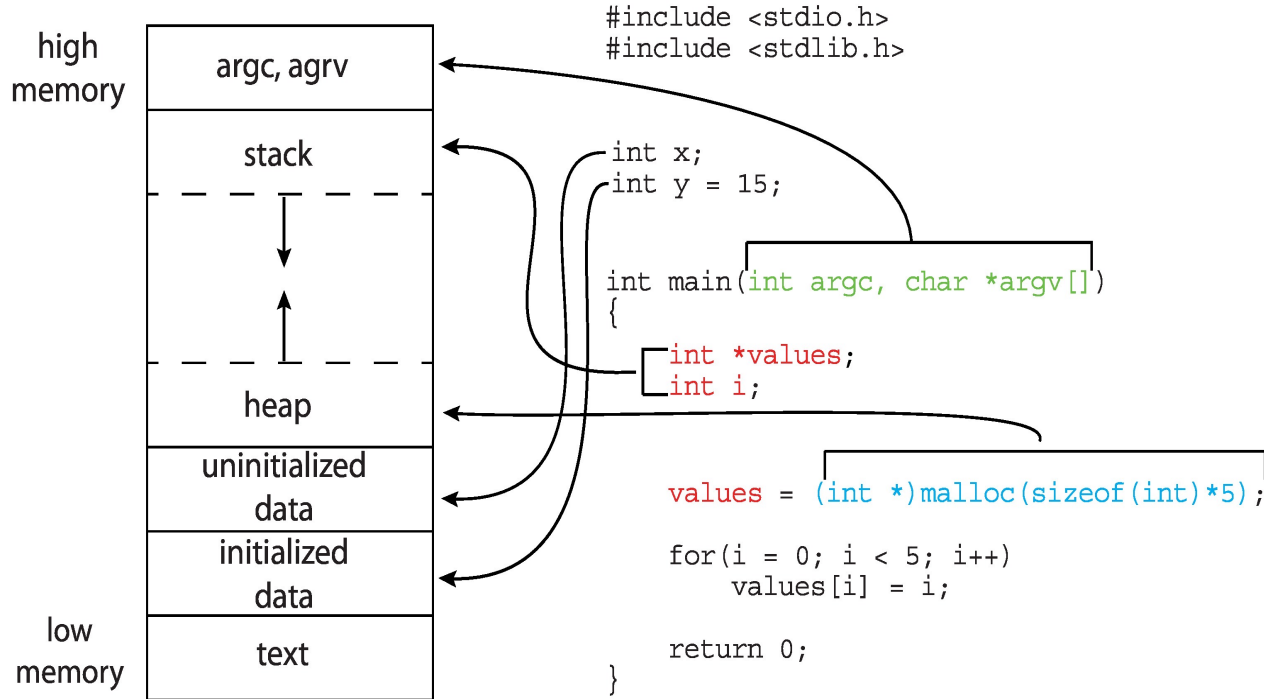
- An operating system executes a variety of programs that run as processes
- **Process** – a program in execution; process execution must progress in sequential fashion
- Multiple parts
 - The **program code**, also called **text section**
 - Current activity including **program counter**, and **processor registers**
 - **Stack section** containing temporary data
 - Function parameters, return addresses, local variables
 - **Data section** containing global variables
 - **Heap section** containing memory dynamically allocated during run time

Process Concept (Cont.)

- *Program* is *passive* entity stored on disk (e.g., *executable file*)
- *Process* is *active* entity
 - Program becomes process when executable file loaded into memory
- *Execution of program* can be started via GUI mouse clicks, command line (CLI) entry of its name, etc.
- One program can be several processes
 - E.g., Consider multiple users executing the same program



Memory Layout of a C Program



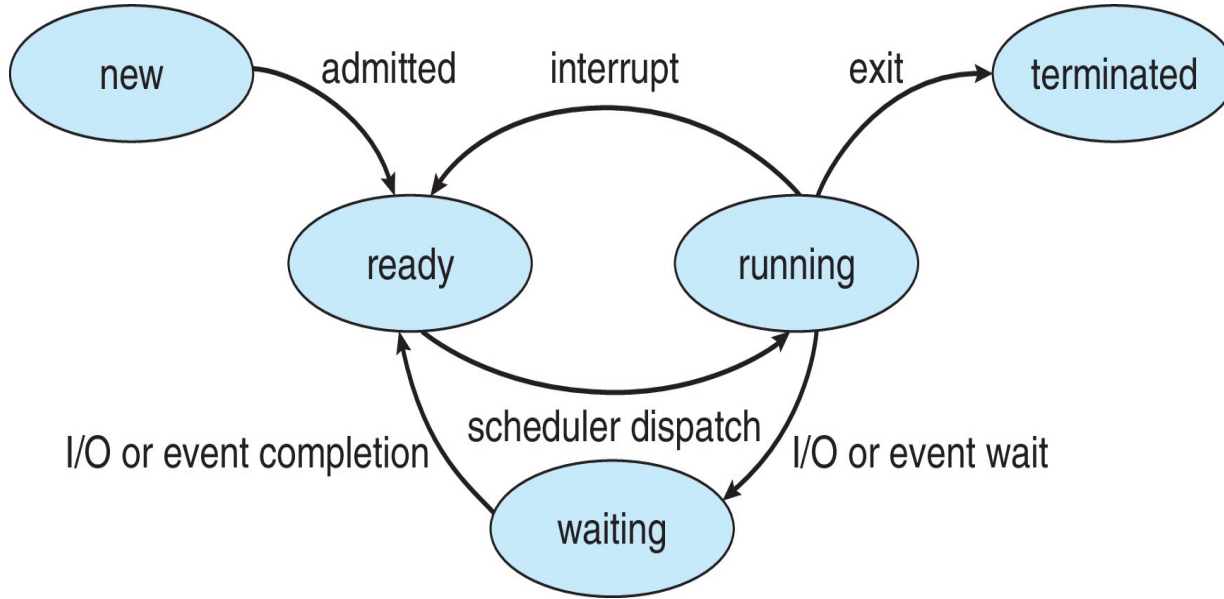
Determine the text/BSS/data/stack segment?

```
#include <iostream>
#include <unistd.h>
int glo_init_data = 9;
int glo_noninit_data;
void print_func () {
    int local_data = 9 ;
    printf( "Process ID = %d\n", getpid());
    printf( "Addresses of the process : \n");
    printf( "  1 . glo_init_data = %p\n", &glo_init_data);
    printf( "  2 . glo_noninit_data = %p\n", &glo_noninit_data);
    printf( "  3 . print_func() = %p\n", &print_func);
    printf( "  4 . local_data = %p\n", &local_data);
}
int main( int argc , char **argv ) {
    print_func();
    return 0;
}
```

Process State

- As a process executes, it changes *state*
 - *New* – The process is being created
 - *Running* – Instructions are being executed
 - *Waiting* – The process is waiting for some event to occur
 - *Ready* – The process is waiting to be assigned to a processor
 - *Terminated* – The process has finished execution

Diagram of Process State



Process Control Block (PCB)

Process Control Block (PCB) – Information associated with each process, also called **Task Control Block (TCB)**, includes:

- *Process state* – running, waiting, etc.
- *Process number* – identity of the process
- *Program counter* – location of instruction to next execute
- *CPU registers* – contents of all process-centric registers
- *CPU scheduling info* – priorities, scheduling queue pointers
- *Memory-management information* – memory allocated to the process
- *Accounting information* – CPU used, clock time elapsed since start, time limits
- *I/O status information* – I/O devices allocated to process, list of open files

process state
process number
program counter
registers
memory limits
list of open files
...

Process IDs

- A unique process ID assigned to each new process by OS
- All process has its parent process (except [the init process](#) , PID=1)

```
#include <stdio.h>
#include <unistd.h>
```

```
int main ( ){
    printf ("The process ID is %d\n", (int) getpid ());
    printf ("The parent process ID is %d\n", (int) getppid ());
    sleep(100);
    return 0 ;
}
```

```
Vans-MacBook-Air:c-examp ltvan$ ./pid
The process ID is 6790
The parent process ID is 6579
```

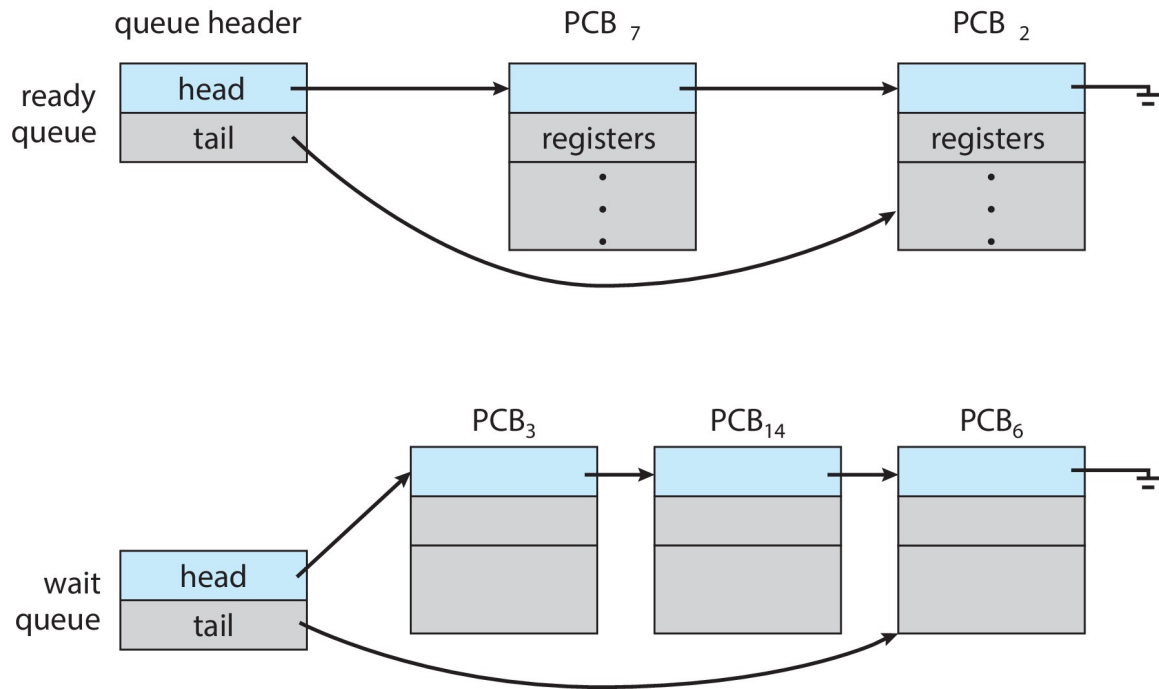
Use [htop](#) command to monitor processes and resources of Linux system

6579	ltvan	32	0	4196M	1240	?	0.0	0.0	0:00.00	└─ -bash
6714	ltvan	25	0	4168M	696	?	0.0	0.0	0:00.00	└─ ./pid

Process Scheduling

- Maximize CPU use → quickly switch processes onto CPU core
- *Process scheduler* selects one process among available (ready) processes for next execution on CPU core
- Maintains *scheduling queues* of processes
 - *Ready queue* – set of all processes residing in main memory, ready and waiting to execute
 - *Wait queues* – set of processes waiting for an event (e.g., I/O)
- Processes migrate among the various queues

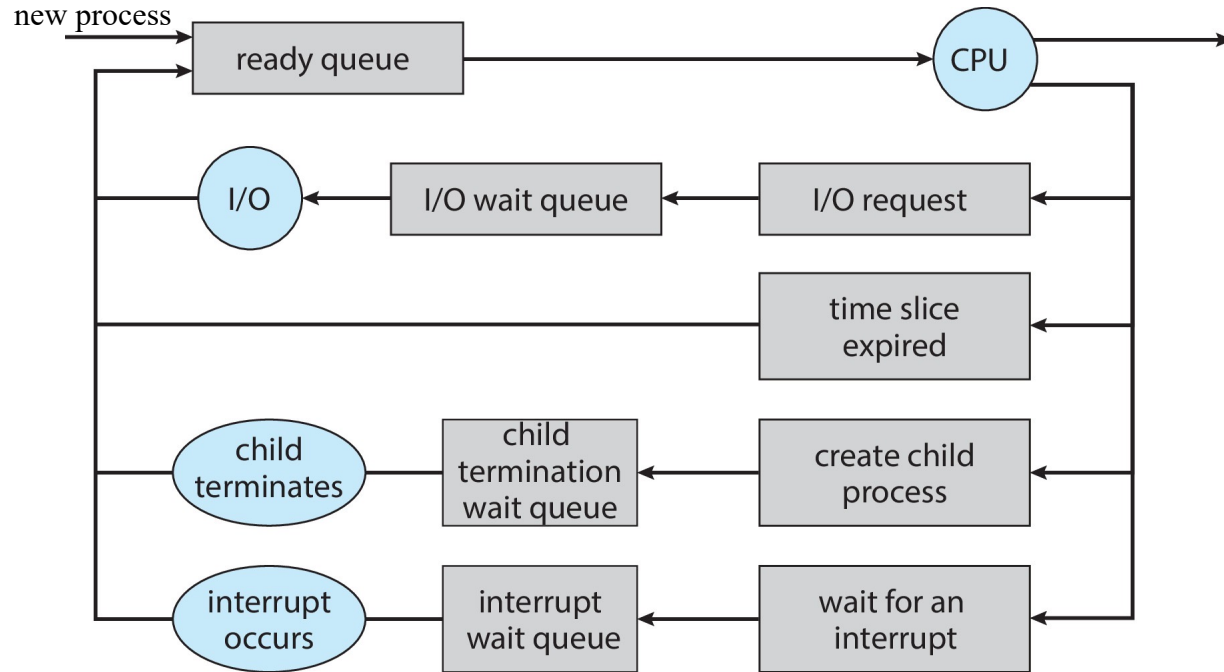
Ready and Wait Queues



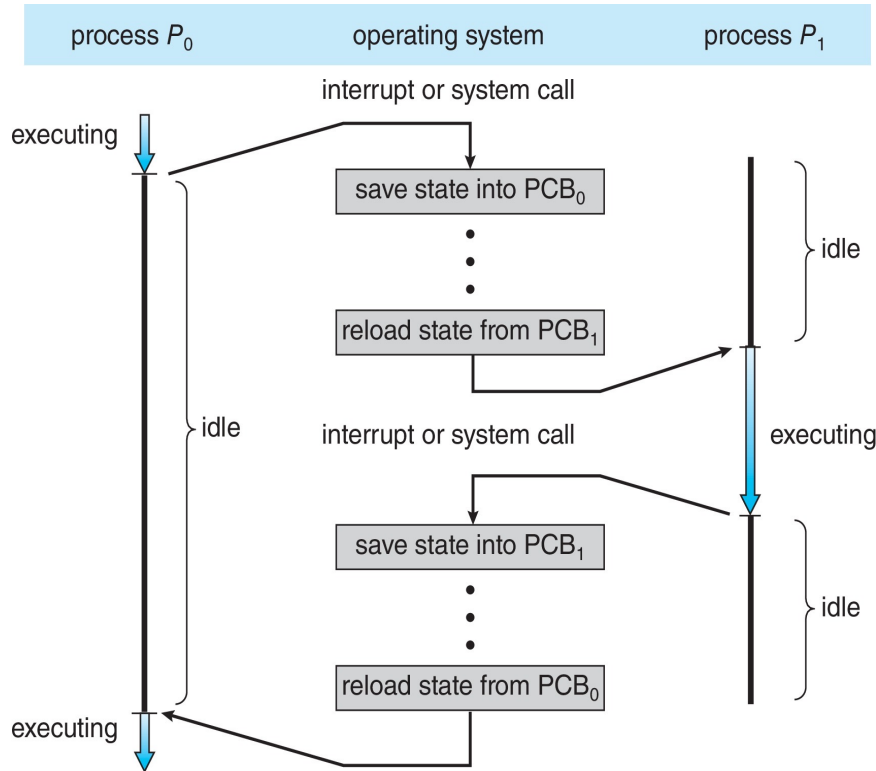
Value “niceness” in process scheduling

- Represent the priority of a process ranged from -20 to +19, 0 by default.
- A higher niceness value means that the process is given a lesser execution priority
- Use **nice command** to assign the niceness of a new process before running:
 - `$sudo nice -n <value> [other arguments] <command>`
- A positive value of niceness reduces the process's execution priority.
- Use **renice command** to change the niceness of a running process from the command line
 - `$sudo renice -n <value> [other arguments] -p <PID>`

Representation of Process Scheduling



CPU Switch from Process to Process



- A *context switch* occurs when the CPU switches from one process to another.

Context Switch

- When CPU switches to another process, the system must *save the state* of the old process and load the *saved state* for the new process via a *context switch*
- *Context* of a process represented in the **PCB**
- Context-switch time is *overhead*, the system does no useful work while switching
 - The more complex the OS and the PCB, the longer the context switch
- *Time* dependent on hardware support
 - Some hardware provides *multiple sets of registers per CPU*, multiple contexts loaded at once

Operations on Processes

- System must provide mechanisms for:
 - process creation
 - process termination

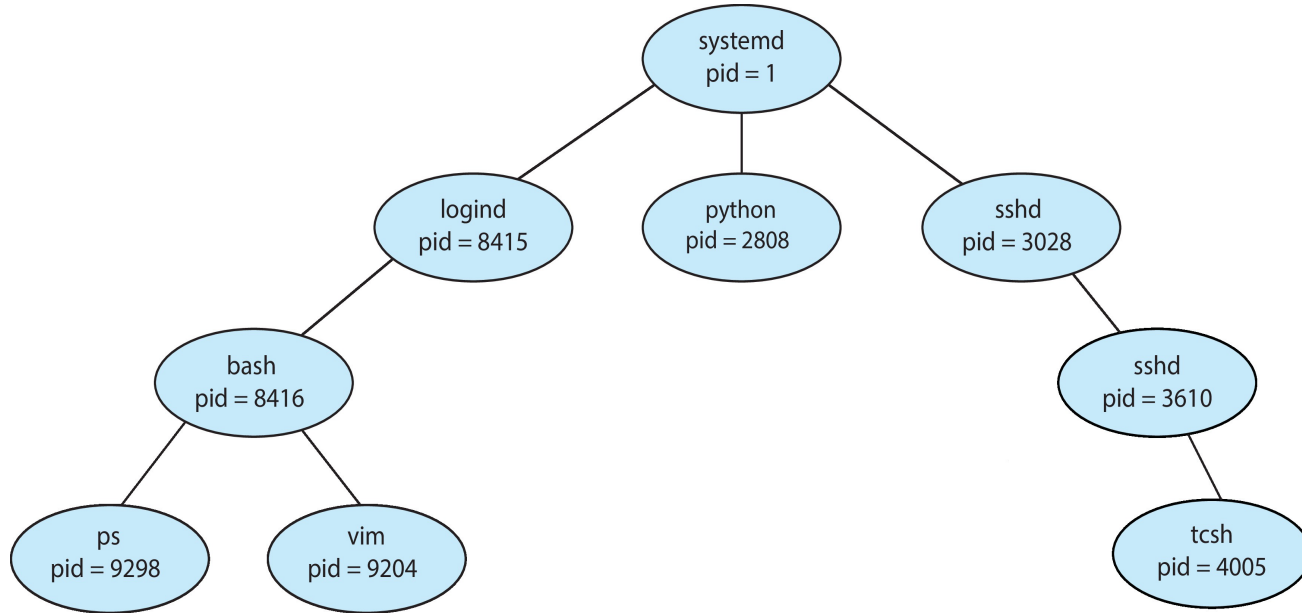
Process Creation

- *Parent processes* create *children processes*, which, in turn create other processes, forming a *tree of processes*
- Process identified and managed via a **Process Identifier (PID)**
- **Resource sharing options**
 - Parent and children share *all* resources
 - Children share *subset* of parent's resources
 - Parent and child share *no* resources

Process Creation (Cont.)

- **Execution options**
 - Parent and children execute concurrently
 - Parent waits until children terminate
- **Address space**
 - Child duplicate of parent
 - Child has a program loaded into it


A Tree of Processes in Linux



#pstree

Process Creation (Cont.)

- UNIX examples
 - **system()** to execute a command from within a program
 - **fork()** system call creates new process
 - **exec()** system call used after a **fork()** to replace the process' memory space with a new program
 - Parent process calls **wait()** waiting for the child to terminate



system() system call

```
#include <stdlib.h>
int system(const char *command);
```

- hands the argument command to the command interpreter sh: `$/bin/sh -c <command>`
 - Ex: `$/bin/sh -c ls`
- If command is **NULL**, system() return a nonzero value if a shell is available, or 0 if no shell is available
- If system() return **-1**: child process could not be created or its status could not be retrieve

```
#include <stdlib.h>
int main ( )
{
    int return_value ;
    return_value = system ( "ls ." );
    return return_value;
}
```

```
Vans-MacBook-Air:c-examp ltvan$ ./systemsimple
2.2.c.rtf      ex             fork2.c        pid
addrspace     ex.c           forkp          pid.cpp
```

exec() system call

- Replaces the current process image with a new process image.
- Can be used to run a C/C++ prog by using another C/C++ prog.
- exec() family:
- **execl(), execlp(), execl()**
 - `int execl(const char *path, const char *arg, ..., NULL);`
 - `int execlp(const char *file, const char *arg, ..., NULL);`
- **execv(), execvp(), execvpe()**
 - `int execv(const char *path, char *const argv[]);`
- **execle(), execvpe()**
 - `int execle(const char *path, const char *arg, ..., NULL, char *const envp[]);`
- **execlp(), execvp(), execvpe()**

C Program Forking A Separate Process: fork() system call

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```

Example

-
- `int main(){`
- `printf("Hello \n");`
- 1. `fork();`
- 2. `fork();`
- `printf("Hello \n");`
- `return 0;`
- `}`

How many processes will be generated?

Process Termination

- Process executes *last statement* and then asks the operating system to delete it using the **exit()** system call.
 - Returns status data from child to parent (via **wait()**)
 - Process' resources are deallocated by operating system
- Parent may terminate the execution of children processes using the **abort()** system call. Some reasons for doing so:
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
 - The parent is exiting and the *operating systems does not allow a child to continue if its parent terminates*

Process Termination (Cont.)

- Some operating systems do not allow child to exist if its parent has terminated. *If a process terminates, then all its children must also be terminated.*
 - **Cascading termination:** All children, grandchildren, etc. are terminated
 - The termination is initiated by the operating system
- The parent process may wait for termination of a child process by using the `wait()` system call. The call returns status information and the **pid** of the terminated process
 - ✓ `pid = wait(&status);`
 - ✓ `waitpid()` suspends execution of the calling process until a child specified by *pid* argument has changed state
- If no parent waiting (did not invoke `wait()`), process is a *zombie*
- If parent terminated without invoking `wait()`, process is an *orphan*

wait(&status) example

```
#include <stdio.h>
#include <unistd.h>
#include <sys/wait.h>
int main (int argc, char *argv[]){
    int    return_code;
    int status;
    /* create a new process */
    return_code = fork();
    if (return_code > 0){
```

```
#include <sys/types.h>
#include <sys/wait.h>
pid_t wait(int *status);
```

```
        int id = wait(&status);
        printf("Parent process: child's pid = %d with returned value = %d \n", id, WEXITSTATUS(status));
        printf("child's pid = %d with raw returned value = %d \n", id, status);
        printf("Pid of child %d \n", return_code);
        return 0;
    }
    else if (return_code == 0){
        printf("This is child process \n");
        return 10; //exit(10);
    }
    else {
        printf("Fork error\n");
        return 1;
    }
}
```

waitpid() example

```
#include <sys/types.h>
```

```
#include <sys/wait.h>
```

```
pid_t waitpid(pid_t pid, int  
*status, int options);
```

- pid < -1: wait for any child process whose process group ID is equal to the absolute value of *pid*.
- pid = 1: wait for any child process
- pid = 0: wait for any child process whose process group ID is equal to that of the calling process
- pid > 0: wait for the child whose process ID is equal to the value of *pid*

```
#include<stdio.h>
#include<sys/wait.h>
#include<unistd.h>

int main(){
    pid_t pid=fork();
    if(pid==0){
        printf("Child proc id = %d, groupid = %d \n",
getpid(),getpgrp());
    }
    else if(pid>0){
        waitpid(pid,NULL,0);
        printf("Parent proc id = %d \n",getpid());
    }
    return 0;
}
```

`wait(&status)` is equivalent to `waitpid(-1, &status, 0);`

```
Vans-MacBook-Air:c-examp ltvan$ ./forkwpidr
Child proc id = 9716, groupid = 9715
Parent proc id = 9715
```

Signals


- A special message sent to a process
- Be processed immediately without finishing the current function or even the current line of code.
- Be specified by a signal number or name.
- A default disposition of each signal determines what happens to the process if the program does not specify some other behavior.
- A process may ignore the signal or call a special signal-handler function to respond to the signal.
- **SIGBUS, SIGSEGV, SIGFPE**: send by Linux system -> terminate the process
- **SIGUSR1, SIGUSR2**: “user-defined” signal send a command to a running program
- **SIGHUP**: wake up an idling program or cause a program to reread its config files.
- **SIGINT**: interrupt from keyboard

Example of signal handling

```
#include<stdio>
#include<signal>
#include<unistd.h>
void sig_handler(int signum){
    //Return type of the handler function should be void
    printf("\nInside handler function\n");
}
int main(){
    signal(SIGINT,sig_handler); // Register signal handler
    for(int i=1;i<10;i++){      //Infinite loop
        printf("%d : Inside main function\n",i);
        sleep(1); // Delay for 1 second
    }
    return 0;
}
```


Multiprocess Architecture – Chrome Browser

- Many web browsers ran as a single process (some still do)
 - If one web site causes trouble, entire browser can hang or crash
- Google Chrome Browser is multiprocess with 3 different types of processes:
 - *Browser process* manages user interface, disk and network I/O
 - *Renderer process* renders web pages, deals with HTML, JavaScript. A new renderer created for each website opened
 - Runs in *sandbox* restricting disk and network I/O, minimizing effect of security exploits
 - *Plug-in process* for each type of plug-in



Each tab represents a separate process.



THANK YOU !

Center Of Computer Engineering